TRAINING AND SCALING LANGUAGE MODELS: FROM FIRST PRINCIPLES TO EFFICIENT SERVING

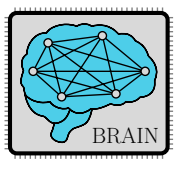
Welcome to Training & Scaling Language Models

Course Presentation: What We Will Do, Organization, and Projects

BRAIN team

IMT Atlantique · January 2027





What can today's language models do?

Read and transform

Summarize, translate, compare, and extract information from text.

Work with code

Generate, explain, debug, and modify programs.

Write and explain

Draft, revise, answer questions, and adapt to instructions or examples.

Solve structured tasks

Follow multi-step instructions and use examples to tackle new tasks.

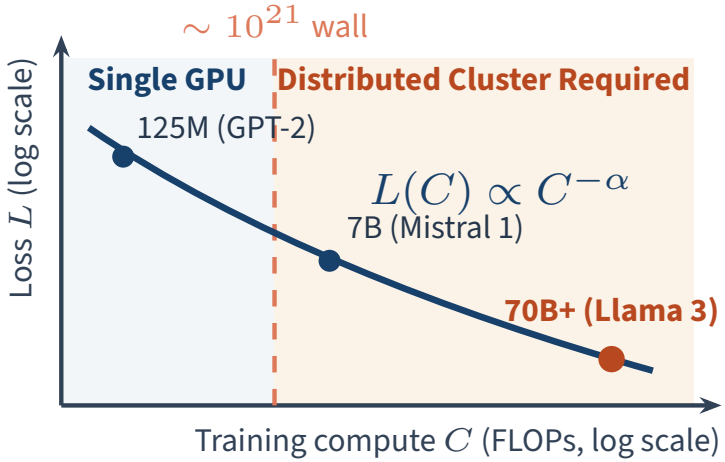
A striking change

Early GPT-2-era models were often limited to short, simple continuations.
Today's models can write code and handle complex tasks.
What drove this performance gap?

Why do scaling laws matter? The empirical mandate of scale

Scaling laws (Compute laws)

- **Empirical power law**: Test loss drops predictably as training compute grows: $\mathbf{L(C) \propto C^{-\alpha_C}}$
- **Predictable returns**: More FLOPs buys lower perplexity, reasoning, and benchmark accuracy.
- **Scaled performance**: Scaling compute unlocks capabilities smaller models cannot reach (e.g. 125M GPT-2 $\rightarrow$ 7B Mistral 1 $\rightarrow$ 70B+ Llama 3).



Frontier scale ($\sim 10^{25}$ FLOPs):

$1 \times$ H100: > 800 **years**

$\downarrow$

$16\text{k} \times$ H100: ≈ 18 **days**

Let's scale models!

What makes large language models hard to build?

As models grow, three limits shape what we can train and run:

1. Compute

GPUs must perform an enormous number of calculations.

2. Memory

The model and its working data must fit in GPU memory.

3. Sharing work

When GPUs share the work, they need to exchange data quickly.

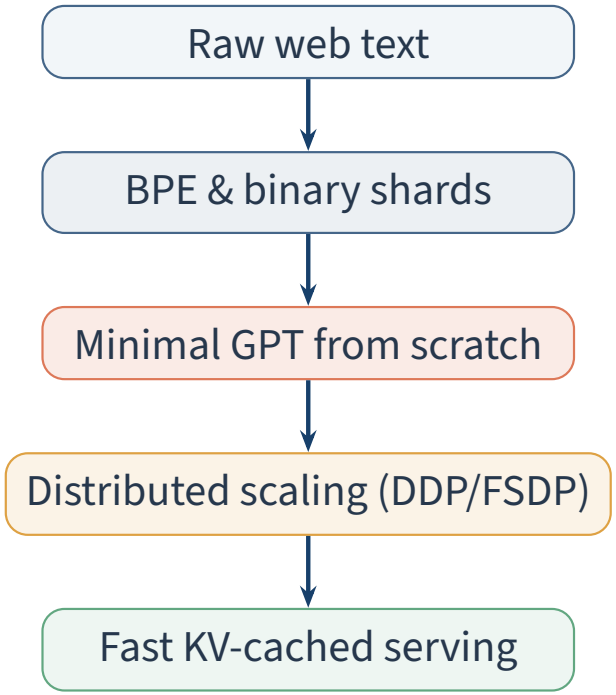
The question throughout this course

How can we train and run useful language models with the time, memory, and hardware we have?

What will you build in this course?

We open the engine from first principles:

- Start with raw, noisy web text.
- Build a decoder-only Transformer from scratch.
- Train, profile, and push a single GPU to its limits.
- Scale across multiple GPUs and nodes.
- Serve the trained model with low latency.



The 3 phases of the course

40 contact hours organized into 32 sessions of 1h15:

Phase 1: Build (15h)

Single-GPU foundation

- Transformer architecture
- Training loop & autograd
- Data pipeline & BPE
- Profiling, MFU, FlashAttn

Phase 2: Scale (15h)

Distributed systems

- Mid-course paper talks
- DDP (all-reduce)
- FSDP & ZeRO sharding
- Tensor & 5D parallelism

Phase 3: Serve (10h)

Inference & synthesis

- KV-cache decode
- Continuous batching
- Open slot (to be decided)
- Final project defense

Each topic connects theory directly to measurable code.

How sessions work: theory + hands-on practice

Theory & derivations

We unpack the mechanics: tensor shapes, autograd graphs, collective communication costs, and roofline models.

Guided implementation

You immediately build, test, and profile the mechanism using our companion PyTorch repository.

No waiting on progress bars

- Class time is for **building, diagnosing, and measuring**.
- Long-running training jobs run asynchronously on the GPU cluster between sessions.
- Next class begins by analyzing the checkpoints, curves, and profiler traces.

A shared course thread, one team research project

Continuous laboratory

The shared thread through the course

- The shared codebase grows with each module.
- Practice building, testing, and measuring.
- Keep working components and evidence as you go.

Reproduction project

Your team's research investigation (65%)

- Choose a paper and test **one precise claim** at smaller scale.
- Reuse or adapt course tools where they help.
- Worth 65%; developed during 4 in-class studio sessions.

Connection: The lab builds shared tools your team can use to test a research claim.

Reproduction project: what we expect

What NOT to do

Do not try to retrain a 70B model or reproduce an entire 20-page paper.

Valid reproduction approaches

- Test a memory or throughput claim under synthetic workloads.
- Test an architectural tweak on a smaller model (e.g. 124M).
- Verify a scaling law or ablation trend under a fixed compute budget.

Negative results welcome

A rigorous experiment that explains *why* a published claim did not hold under controlled settings gets full credit.

Scientific rigor

Controlled variables, verified baselines, and clear ablations matter far more than raw model size.

Grading and assessment

We grade engineering rigor, correct implementations, and controlled reasoning:

Component	Weight	Format
Pre-session quizzes	**15%**	5-min online checks before each teaching day; verifies reading & core concepts
Mid-course paper talk	**20%**	8-min presentation + 4-min Q&A; explain the paper's mechanism & proposed test
Reproduction project	**65%**	Reproducible code, short technical report ($\leq$ 5 pages), and team defense with individual Q&A

Quizzes drop the lowest scores to accommodate occasional technical issues or absences.

Course toolkit and logistics

- **Companion repo**: `llm-course-companion`
- Clean, modular PyTorch implementation for each module.
- Automated test suites:
  - Tensor shape invariants
  - Causal masking checks
  - Tiny-batch overfit ($L < 0.1$)
  - Checkpoint resume determinism
- Fast dependency management with `uv`.
- Dedicated GPU compute nodes for training runs.

Sanity check before Session 1

```
# Clone & install
git clone <repo-url>
cd llm-course-companion
uv sync

# Check GPU & tests
nvidia-smi
uv run pytest tests/
```

Summary

1. **First principles**: No black boxes. We understand every tensor shape, CUDA kernel launch, and collective communication.
2. **Measure before optimizing**: Trust profiler traces and GPU metrics, not intuition.
3. **Build together, explore in teams**: Master the fundamentals in the continuous lab, then test cutting-edge research in your reproduction project.
4. **Rigorous science**: Reproducible baselines and honest analysis of trade-offs.

Let’s build a language model!