

Session 11: KV-Cached Decoding

Reuse, Memory Growth, and Exact Autoregressive State

Jonathan Lys

IMT Atlantique
September 21, 2026



Training Language Models: From First Principles to Efficient Serving

Lecture map

- ① From full-prefix recomputation to reuse
- ② Cache memory and bandwidth
- ③ Correctness and benchmarking

Guiding question

What can autoregressive decoding reuse exactly, and how does that reuse turn compute into a memory-management problem?

What should be clear by the end?

- ① Explain the redundant work removed by a KV cache.
- ② Compute cache shapes and memory for MHA, GQA, and MQA.
- ③ Reason about decode bandwidth and cache layout.
- ④ Test cached generation against a full-prefix reference.

Outline

- 1 From full-prefix recomputation to reuse
- 2 Cache memory and bandwidth
- 3 Correctness and benchmarking

Naive decoding recomputes the unchanged prefix

To generate token $t + 1$, a naive decoder runs the model on $x_{0:t}$. The next step runs it again on $x_{0:t+1}$.

Repeated work

At every layer, keys and values for positions $0 \dots t$ are deterministic functions of unchanged hidden states. Recomputing them cannot change the result.

Cache principle

Store past key and value projections once. At the next step, compute projections only for the new token and attend to cached history.

Source: Saad Ahmed, KV Cache Explained Intuitively

Prefill and decode have different shapes

Prefill

processes the prompt in parallel
many query tokens
large matrix multiplications
usually compute-efficient


Decode

processes one new token per request
reads growing KV history
small matrix operations
often bandwidth-bound

Do not average them together

Time to first token and inter-token latency arise from different workloads and need separate measurements.

Cached attention at one layer



$$o_t = \text{softmax}\left(q_t K_{1:t}^\top / \sqrt{d_h}\right) V_{1:t}$$

The cache removes repeated K, V projection; attention still scans the growing context.

The cache changes asymptotic projection work

For n generated tokens after a prompt:

Without cache

Reproject the growing prefix at every step. Prefix projection work accumulates roughly quadratically in generated length.

With cache

Project each token's key and value once. Attention still scans a context that grows linearly per decode step.

The cache removes redundant projection work. It does not make full-context attention constant-time.

Outline

- ① From full-prefix recomputation to reuse
- ② Cache memory and bandwidth
- ③ Correctness and benchmarking

Cache shape determines capacity

For L layers, batch B , H_{kv} key-value heads, cached length T , head width d_h , and element size s :

$$M_{KV} = 2 L B H_{kv} T d_h s$$

The factor 2 stores both keys and values.

Per-token, per-request growth

$$m_{\text{token}} = 2 L H_{kv} d_h s$$

This number links model architecture directly to serving concurrency.

Worked memory estimate

Suppose $L = 32$, $H_{kv} = 32$, $d_h = 128$, BF16, and $T = 4096$:

$$M_{KV} = 2 \cdot 32 \cdot 1 \cdot 32 \cdot 4096 \cdot 128 \cdot 2 \approx 2 \text{ GiB}$$

Serving implication

Sixteen concurrent requests at this cached length would need roughly 32 GiB for KV state alone, before weights, workspaces, and fragmentation.

MHA, GQA, and MQA differ mainly in H_{kv}

Variant	query heads	KV heads	cache effect
MHA	H	H	baseline
GQA	H	G , where $1 < G < H$	reduced by H/G
MQA	H	1	reduced by H

Query heads share key-value groups while retaining distinct queries. This reduces cache capacity and decode bandwidth at a possible modeling-quality trade-off.

Decode often reads far more bytes than it computes

Each new token reads the cached K , V for every active layer and request. Batch-1 decode has little reuse across queries.

Bandwidth model

$$t_{\text{decode}} \gtrsim \frac{\text{weight bytes read} + \text{KV bytes read}}{BW_{\text{effective}}}$$

Consequence

Reducing KV heads, using lower cache precision, batching requests, and improving memory layout can matter more than reducing arithmetic.

Allocation strategy decides fragmentation and copy cost

Concatenate

- simple
- repeated allocation and copying
- poor for long generation

Preallocate

- fast append
- reserves maximum length
- internal waste

Paged

- non-contiguous blocks
- flexible growth
- mapping metadata

Paged layouts make variable-length request lifetimes manageable without requiring one contiguous maximum-size allocation per request.

Checkpoint: capacity calculation

Think, then compare

Recompute the worked example with eight KV heads. How many 4096-token requests fit in 32 GiB if 24 GiB is available for cache after weights and workspaces?

Outline

- ① From full-prefix recomputation to reuse
- ② Cache memory and bandwidth
- ③ Correctness and benchmarking

The cache must preserve logical positions

- The new token uses its absolute or rotary position t .
- The causal mask covers the valid cached prefix and current token.
- Padding and per-request lengths cannot expose uninitialized cache entries.
- Reordered or paged blocks must reconstruct the same logical sequence.

Common failure

Correct key-value values at the wrong positions produce plausible but divergent text.

Full-prefix decoding is the reference oracle

For fixed weights, prompt, and token sequence:

- ① run a full-prefix forward pass for each step
- ② run cached prefill plus single-token decode
- ③ compare next-token logits at every step
- ④ repeat across batch sizes, prompt lengths, and cache boundaries

Then test generation

Greedy outputs should match exactly within numerical tolerance. Sampling tests require identical random state and probability inputs.

Benchmark prefill and decode separately

Phase	Primary metrics	Sweep
prefill	latency, prompt tokens/s, peak memory	prompt length, batch
decode	ms/token, output tokens/s, cache bytes	cached length, batch
end-to-end	TTFT, total latency	prompt/output distribution

Warm and synchronize

Kernel compilation, allocation, and asynchronous timing can otherwise dominate or under-report short decode steps.

Checkpoint: cache debugging

Think, then compare

Cached logits match for short prompts but diverge immediately after a page boundary. Which state and indexing checks isolate layout bugs from positional-encoding bugs?

The durable ideas

- ① KV caching reuses past key and value projections exactly.
- ② Cache memory grows with layers, KV heads, requests, cached tokens, head width, and precision.
- ③ Decode is commonly bandwidth-bound, so cache architecture and layout shape performance.
- ④ Full-prefix logit equivalence is the primary correctness test before benchmarking.